

TCP EX MACHINA: COMPUTER-GENERATED CONGESTION CONTROL

KEITH WINSTEIN AND HARI BALAKRISHNAN

Presented by: Angela Jiang

Network congestion

- Cause: Sources trying to send data faster than the network can process
- Result: QoS deterioration of network
 - Queuing delays
 - Packet losses
 - Congestion collapse

Congestion control

- Prevent congestion collapse
- Allocates network resources
 - Link bandwidth
 - Queue space
- End-to-end or network assisted

End-to-end congestion control

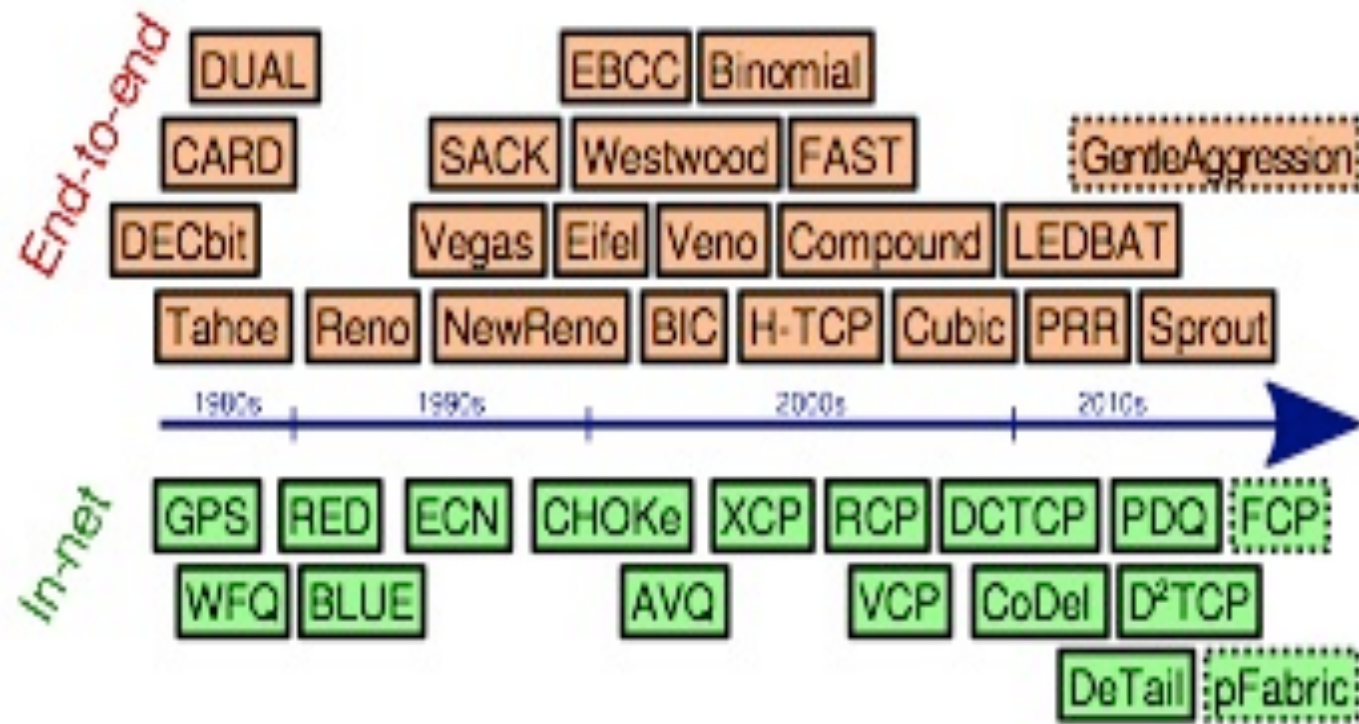
- No support from network layer
- Hosts must infer presence of congestion
 - Packet loss
 - Queueing delay
 - ECN marks
- Individually control transmission rate

Challenge: evolving networks



- Wireless
- Short connections
- Bursty traffic
- Datacenters

Evolution of congestion control



Challenge: rational choice of scheme

Cubic vs. Compound

- Different goals?
- Different assumptions about network?
- One scheme just plain better?

Current congestion control

- Inflexible; doesn't allow network evolution
- Unclear what an algorithm is optimized for

Remy

- *Program that generates end-to-end congestion control schemes offline*
- **Given**
 - Network representation
 - Objective of app (e.g. high throughput)
- **Generates**
 - RemyCC; a congestion control algorithm

Objective function

- Fairness vs. efficiency

$$U_{\alpha}(x) = \frac{x^{1-\alpha}}{1-\alpha}$$

- Delay vs. throughput

$$U_{\alpha}(x) - \delta \cdot U_{\beta}(y)$$

- Objectives used

$$U = \log(\text{throughput}) - \delta \cdot \log(\text{delay})$$

$$U = -\frac{1}{\text{throughput}}$$

Prior assumptions of network

- Model of network uncertainty
 - Link speed distribution
 - Delay distribution
 - Degree of multiplexing
- Traffic model
 - Off-to-on model
 - Web browsing, MapReduce, VoIP

RemyCC maps state to an action

$$\text{RULE}(r_ewma, s_ewma, rtt_ratio) \rightarrow \langle m, b, \tau \rangle$$

- m Multiple to congestion window
- b Increment to congestion window
- τ Minimum interval between two outgoing packets

Find the best value

r_ewma


<?, ?, ?>

s_ewma

Best single action




Subdivide most used rule

r_ewma	$\langle 0.90, 4, 3.3 \rangle$	$\langle 0.90, 4, 3.3 \rangle$
	$\langle 0.90, 4, 3.3 \rangle$	$\langle 0.90, 4, 3.3 \rangle$
s_ewma		

Optimize each new action

r_ewma	<0.90,5,2.8>	<0.70,6,53.5>
	<0.60,19,76.2>	<0.80,5,4.1>
		s_ewma


Split most used rule



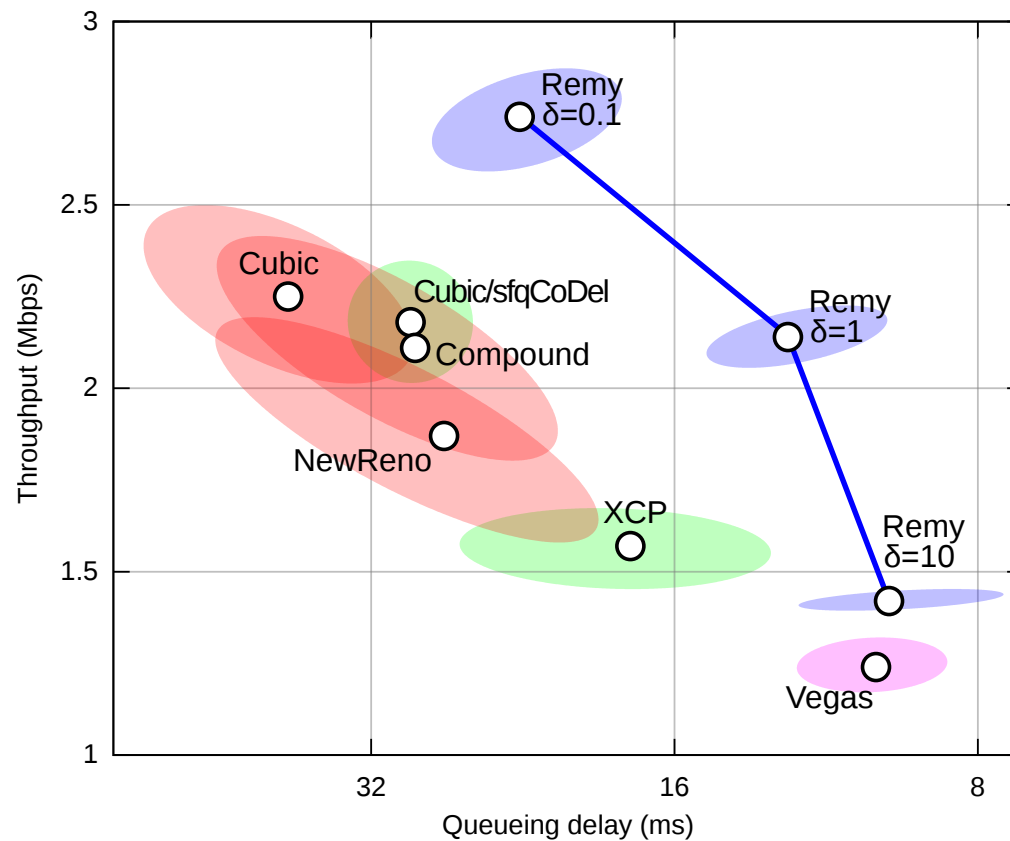
Evaluation

- 3 RemyCCs with $\delta = .1, 1, 10$
- Generating one RemyCC
 - Takes a few hours
 - \$5-\$10 on EC2
- Compared against end-to-end and network assisted congestion control schemes


Single bottleneck (“dumbbell”)




Varying throughput (“cellular”)



Cellular: n=8



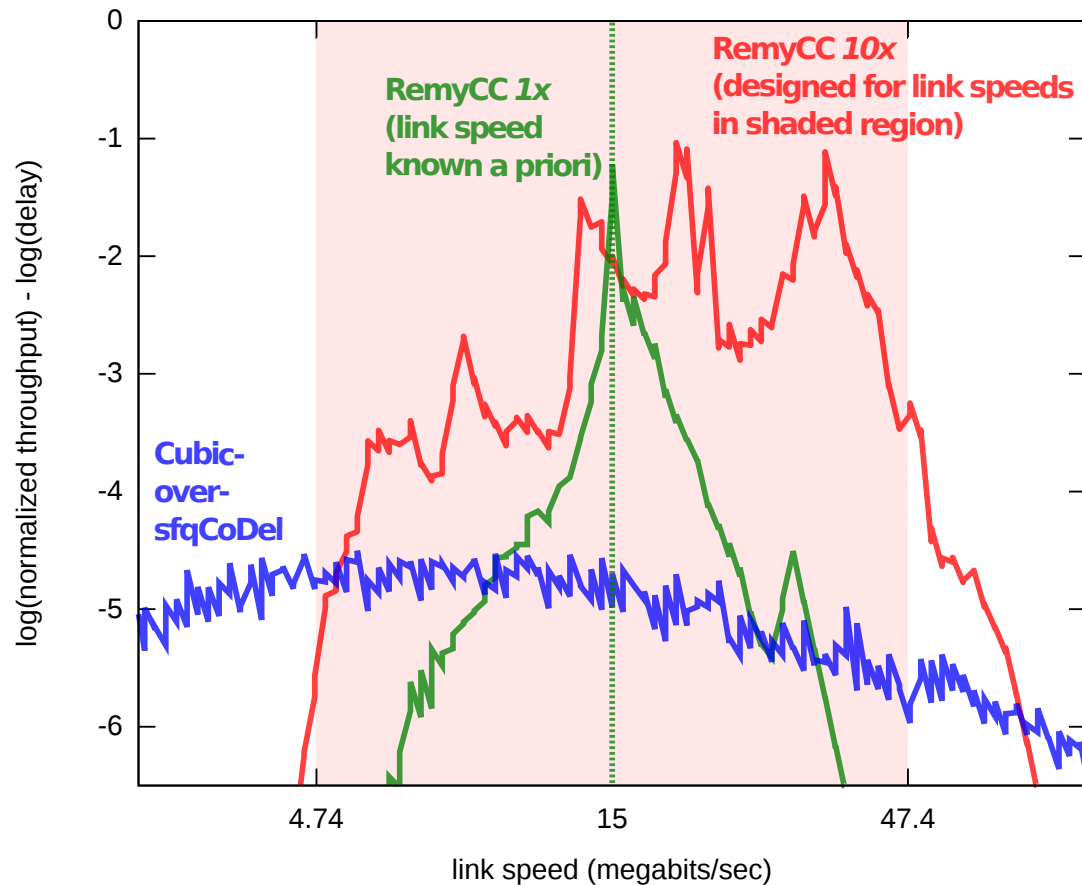
Fairness for varying RTTs



Datacenters

- Issues
 - Many synchronous requests -> incast
 - Diverse mix of short and long flows
- Simulation
 - 64 connections, 10 Gpbs link
 - Objective function maximizes throughput
 - Comparable throughputs to DCTCP
 - DCTCP has shorter RTTs

Sensitivity of design range



Discussion: practicality

- Providing network assumptions
- Scalability
- Using CC algorithms we don't understand
- Coexisting with other protocols

Approaching congestion control differently

- Complex rules but consistent behavior
- Objective and environment driven
- Able to evolve